

# dm-1-asst-3

November 25, 2023

```
[1]: import pandas as pd
import seaborn as sns
import os
import keras
import matplotlib.pyplot as plt
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Conv2D, Dense, MaxPool2D, Flatten, Dropout, \
    BatchNormalization, DepthwiseConv2D, InputLayer, GlobalAveragePooling2D
from tensorflow.keras.regularizers import l2
from tensorflow.keras.models import Sequential
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

## 1. Association Rule Generation from Transaction Data

```
[3]: dataset = pd.read_csv('/content/drive/MyDrive/Grocery_Items_25.csv')
dataset.head()
```

```
[3]:
```

	0	1	2	\
0	pastry	baking powder	NaN	
1	sausage	sausage	frozen fish	
2	ham	whole milk	NaN	
3	whipped/sour cream	baking powder	NaN	
4	house keeping products	frozen vegetables	NaN	

	3	4	5	6	7	8	9	10
0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	house keeping products	NaN	NaN	NaN	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

```
[4]: dataset.shape
```

```
[4]: (8000, 11)
```

```
[5]: records = []  
for i in range(8000):  
    records.append([str(dataset.values[i,j]) for j in range(11) if str(dataset.  
    ↪values[i,j]) != 'nan'])
```

```
[6]: encoder = TransactionEncoder()  
encoder_array = encoder.fit(records).transform(records)  
dataframe = pd.DataFrame(encoder_array, columns=encoder.columns_)
```

```
[7]: dataframe.head()
```

```
[7]:      Instant food products  UHT-milk  abrasive cleaner  artif. sweetener  bags  \  
0          False          False          False          False  False  False  
1          False          False          False          False  False  False  
2          False          False          False          False  False  False  
3          False          False          False          False  False  False  
4          False          False          False          False  False  False  
  
      baking powder  bathroom cleaner  beef  berries  beverages  ...  turkey  \  
0          True          False  False  False  False  ...  False  
1          False          False  False  False  False  ...  False  
2          False          False  False  False  False  ...  False  
3          True          False  False  False  False  ...  False  
4          False          False  False  False  False  ...  False  
  
      vinegar  waffles  whipped/sour cream  whisky  white bread  white wine  \  
0      False  False          False  False  False  False  
1      False  False          False  False  False  False  
2      False  False          False  False  False  False  
3      False  False          True  False  False  False  
4      False  False          False  False  False  False  
  
      whole milk  yogurt  zwieback  
0      False  False  False  
1      False  False  False  
2      True  False  False  
3      False  False  False  
4      False  False  False  
  
[5 rows x 165 columns]
```

C

```
[8]: frequent_itemsets = apriori(dataframe, min_support=0.01, use_colnames=True)
rules = association_rules(frequent_itemsets, metric="confidence",
    ↪min_threshold=0.1)
rules
```

```
[8]:
```

	antecedents	consequents	antecedent support \
0	(soda)	(other vegetables)	0.103125
1	(other vegetables)	(whole milk)	0.122000
2	(rolls/buns)	(whole milk)	0.109625
3	(soda)	(whole milk)	0.103125
4	(yogurt)	(whole milk)	0.088125

	consequent support	support	confidence	lift	leverage	conviction \
0	0.122000	0.010625	0.103030	0.844511	-0.001956	0.978851
1	0.154625	0.014250	0.116803	0.755397	-0.004614	0.957176
2	0.154625	0.012750	0.116306	0.752178	-0.004201	0.956637
3	0.154625	0.011875	0.115152	0.744715	-0.004071	0.955390
4	0.154625	0.011375	0.129078	0.834781	-0.002251	0.970667

	zhangs_metric
0	-0.170323
1	-0.269433
2	-0.270093
3	-0.276522
4	-0.178338


D

```
[9]: heatmap_data = []
for i in [0.001,0.005,0.01,0.05]:
    for j in [0.05,0.075,0.1]:
        frequent_itemsets = apriori(dataframe, min_support=i, use_colnames=True)
        rules = association_rules(frequent_itemsets, metric="confidence",
    ↪min_threshold=j)
        heatmap_data.append({'msv': i, 'mct': j, 'no.of rules': len(rules)})
```

```
[10]: heatdata = pd.DataFrame(heatmap_data)
heatdata = heatdata.pivot(index='mct', columns='msv', values='no.of rules')
```

```
[11]: sns.heatmap(heatdata.sort_index(ascending=False), cmap='coolwarm')
```

```
[11]: <Axes: xlabel='msv', ylabel='mct'>
```



E

```
[12]: frequent_itemsets_hc = apriori(dataframe, min_support=0.005, use_colnames=True)
rules_hc = association_rules(frequent_itemsets_hc, metric="confidence",
                             min_threshold=0.0)
rules_hc.loc[rules_hc['confidence'].idxmax()]
```

```
[12]: antecedents      (domestic eggs)
consequents      (whole milk)
antecedent support      0.035
consequent support      0.154625
support      0.005625
confidence      0.160714
lift      1.039381
leverage      0.000213
conviction      1.007255
zhangs_metric      0.039263
Name: 15, dtype: object
```

## 2. Image Classification using CNN

```
[13]: directory = "/content/drive/MyDrive/Dogs_Dataset/images_cropped"
```

```
[14]: Image = []
label = []
categories = sorted(os.listdir(directory))
for i in categories:
    for j in os.listdir(os.path.join(directory,i)):
        Image.append(os.path.join(directory,i,j))
        label.append(i)
```

```
[15]: train_data,test_data,train_label,test_label = train_test_split(Image,
                                                                    label,
                                                                    train_size=0.8,
                                                                    random_state=10)
```

```
[16]: train = pd.DataFrame({'file_paths': train_data, 'labels': train_label})
test = pd.DataFrame({'file_paths': test_data, 'labels': test_label})
```

```
[17]: datagen = keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
train_gen = datagen.flow_from_dataframe(
    train,
    x_col='file_paths',
    y_col='labels',
    target_size=(256, 256),
    batch_size=32,
    class_mode='categorical',
)
test_gen = datagen.flow_from_dataframe(
    test,
    x_col='file_paths',
    y_col='labels',
    target_size=(256, 256),
    batch_size=32,
    class_mode='categorical',
)
```

Found 543 validated image filenames belonging to 4 classes.

Found 136 validated image filenames belonging to 4 classes.

```
[18]: model = Sequential()

model.add(Conv2D(8, (3, 3), activation='relu', input_shape=(256,256,3)))
model.add(MaxPool2D((2, 2)))
model.add(Flatten())
model.add(Dense(16, activation='relu'))
model.add(Dense(4, activation='softmax'))
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 8)	224
max_pooling2d (MaxPooling2D)	(None, 127, 127, 8)	0
flatten (Flatten)	(None, 129032)	0
dense (Dense)	(None, 16)	2064528
dense_1 (Dense)	(None, 4)	68

Total params: 2,064,820  
Trainable params: 2,064,820  
Non-trainable params: 0

```
[19]: training = model.fit(train_gen,  
                           validation_data = test_gen,  
                           epochs = 20  
                           )
```

Epoch 1/20

17/17 [=====] - 108s 6s/step - loss: 3.4264 - accuracy: 0.2523 - val\_loss: 1.3911 - val\_accuracy: 0.2574

Epoch 2/20

17/17 [=====] - 22s 1s/step - loss: 1.3826 - accuracy: 0.2394 - val\_loss: 1.3997 - val\_accuracy: 0.2426

Epoch 3/20

17/17 [=====] - 21s 1s/step - loss: 1.3806 - accuracy: 0.2431 - val\_loss: 1.3977 - val\_accuracy: 0.2426

Epoch 4/20

17/17 [=====] - 21s 1s/step - loss: 1.3781 - accuracy: 0.2449 - val\_loss: 1.3985 - val\_accuracy: 0.2500

Epoch 5/20

17/17 [=====] - 22s 1s/step - loss: 1.3749 - accuracy: 0.2210 - val\_loss: 1.3921 - val\_accuracy: 0.2868

Epoch 6/20

17/17 [=====] - 21s 1s/step - loss: 1.3565 - accuracy: 0.2744 - val\_loss: 1.3635 - val\_accuracy: 0.3382

Epoch 7/20

```

17/17 [=====] - 20s 1s/step - loss: 1.2699 - accuracy:
0.4309 - val_loss: 1.2837 - val_accuracy: 0.4044
Epoch 8/20
17/17 [=====] - 21s 1s/step - loss: 1.1688 - accuracy:
0.5562 - val_loss: 1.2000 - val_accuracy: 0.4191
Epoch 9/20
17/17 [=====] - 20s 1s/step - loss: 1.0566 - accuracy:
0.6114 - val_loss: 1.2491 - val_accuracy: 0.3971
Epoch 10/20
17/17 [=====] - 21s 1s/step - loss: 0.9540 - accuracy:
0.6390 - val_loss: 1.2111 - val_accuracy: 0.5662
Epoch 11/20
17/17 [=====] - 21s 1s/step - loss: 0.8853 - accuracy:
0.6796 - val_loss: 1.1892 - val_accuracy: 0.4926
Epoch 12/20
17/17 [=====] - 21s 1s/step - loss: 0.8399 - accuracy:
0.6354 - val_loss: 1.1656 - val_accuracy: 0.5735
Epoch 13/20
17/17 [=====] - 21s 1s/step - loss: 0.8073 - accuracy:
0.6501 - val_loss: 1.1357 - val_accuracy: 0.5882
Epoch 14/20
17/17 [=====] - 21s 1s/step - loss: 0.7715 - accuracy:
0.7182 - val_loss: 1.1178 - val_accuracy: 0.5074
Epoch 15/20
17/17 [=====] - 21s 1s/step - loss: 0.7529 - accuracy:
0.7256 - val_loss: 1.1355 - val_accuracy: 0.5147
Epoch 16/20
17/17 [=====] - 22s 1s/step - loss: 0.7438 - accuracy:
0.7293 - val_loss: 1.0690 - val_accuracy: 0.5515
Epoch 17/20
17/17 [=====] - 23s 1s/step - loss: 0.7329 - accuracy:
0.7403 - val_loss: 1.1131 - val_accuracy: 0.5147
Epoch 18/20
17/17 [=====] - 21s 1s/step - loss: 0.7039 - accuracy:
0.7330 - val_loss: 1.0837 - val_accuracy: 0.5221
Epoch 19/20
17/17 [=====] - 21s 1s/step - loss: 0.6916 - accuracy:
0.7551 - val_loss: 1.1212 - val_accuracy: 0.5074
Epoch 20/20
17/17 [=====] - 21s 1s/step - loss: 0.6631 - accuracy:
0.7680 - val_loss: 1.0856 - val_accuracy: 0.5221

```

```

[20]: def plot_accuracy(train_data):
    acc = train_data.history['accuracy']
    val_acc = train_data.history['val_accuracy']


    epochs = range(len(acc))

```

```

fig = plt.figure(figsize=(10,6))
plt.plot(epochs,acc,c="red",label="Training")
plt.plot(epochs,val_acc,c="blue",label="Validation")
plt.xticks(fontsize=10)
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
plot_accuracy(training)

```



Rowan Banner ID - 916462027

Hence, C) Train the CNN using 2 other number of nodes in the hidden layer 8 and 32 with all other parameters unchanged

```

[22]: model_2 = Sequential()

model_2.add(Conv2D(8, (3, 3), activation='relu', input_shape=(256,256,3)))
model_2.add(MaxPool2D((2, 2)))
model_2.add(Flatten())
model_2.add(Dense(8, activation='relu'))
model_2.add(Dense(4, activation='softmax'))
model_2.compile(optimizer='adam',
                loss='categorical_crossentropy',

```



```

        metrics=['accuracy'])

model_2.summary()

training_2 = model_2.fit(train_gen,
                        validation_data = test_gen,
                        epochs = 20
                        )

```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 254, 254, 8)	224
max_pooling2d_2 (MaxPooling 2D)	(None, 127, 127, 8)	0
flatten_2 (Flatten)	(None, 129032)	0
dense_4 (Dense)	(None, 8)	1032264
dense_5 (Dense)	(None, 4)	36

```

=====
Total params: 1,032,524
Trainable params: 1,032,524
Non-trainable params: 0

```

```

-----
Epoch 1/20
17/17 [=====] - 22s 1s/step - loss: 3.9095 - accuracy:
0.2578 - val_loss: 1.4346 - val_accuracy: 0.3382
Epoch 2/20
17/17 [=====] - 20s 1s/step - loss: 1.3623 - accuracy:
0.3168 - val_loss: 1.3316 - val_accuracy: 0.3529
Epoch 3/20
17/17 [=====] - 28s 2s/step - loss: 1.1848 - accuracy:
0.4125 - val_loss: 1.3236 - val_accuracy: 0.3676
Epoch 4/20
17/17 [=====] - 22s 1s/step - loss: 1.0577 - accuracy:
0.4622 - val_loss: 1.3026 - val_accuracy: 0.4044
Epoch 5/20
17/17 [=====] - 19s 1s/step - loss: 0.9561 - accuracy:
0.4843 - val_loss: 1.3023 - val_accuracy: 0.3824
Epoch 6/20
17/17 [=====] - 21s 1s/step - loss: 0.8640 - accuracy:
0.5009 - val_loss: 1.3390 - val_accuracy: 0.4044


```

```

Epoch 7/20
17/17 [=====] - 22s 1s/step - loss: 0.7994 - accuracy:
0.5046 - val_loss: 1.3293 - val_accuracy: 0.4265
Epoch 8/20
17/17 [=====] - 21s 1s/step - loss: 0.7424 - accuracy:
0.5230 - val_loss: 1.3229 - val_accuracy: 0.3897
Epoch 9/20
17/17 [=====] - 21s 1s/step - loss: 0.7031 - accuracy:
0.5378 - val_loss: 1.3519 - val_accuracy: 0.4338
Epoch 10/20
17/17 [=====] - 23s 1s/step - loss: 0.6647 - accuracy:
0.5506 - val_loss: 1.3860 - val_accuracy: 0.4338
Epoch 11/20
17/17 [=====] - 23s 1s/step - loss: 0.6386 - accuracy:
0.5875 - val_loss: 1.3700 - val_accuracy: 0.4485
Epoch 12/20
17/17 [=====] - 21s 1s/step - loss: 0.6167 - accuracy:
0.6354 - val_loss: 1.4254 - val_accuracy: 0.4559
Epoch 13/20
17/17 [=====] - 22s 1s/step - loss: 0.5899 - accuracy:
0.6556 - val_loss: 1.4445 - val_accuracy: 0.4485
Epoch 14/20
17/17 [=====] - 22s 1s/step - loss: 0.5615 - accuracy:
0.7256 - val_loss: 1.4723 - val_accuracy: 0.4853
Epoch 15/20
17/17 [=====] - 19s 1s/step - loss: 0.5290 - accuracy:
0.8085 - val_loss: 1.6248 - val_accuracy: 0.4265
Epoch 16/20
17/17 [=====] - 20s 1s/step - loss: 0.5032 - accuracy:
0.8453 - val_loss: 1.6332 - val_accuracy: 0.4044
Epoch 17/20
17/17 [=====] - 20s 1s/step - loss: 0.4767 - accuracy:
0.8453 - val_loss: 1.6817 - val_accuracy: 0.3824
Epoch 18/20
17/17 [=====] - 20s 1s/step - loss: 0.4291 - accuracy:
0.9392 - val_loss: 1.7655 - val_accuracy: 0.3824
Epoch 19/20
17/17 [=====] - 21s 1s/step - loss: 0.3984 - accuracy:
0.9429 - val_loss: 1.8532 - val_accuracy: 0.3824
Epoch 20/20
17/17 [=====] - 22s 1s/step - loss: 0.3525 - accuracy:
0.9797 - val_loss: 1.9700 - val_accuracy: 0.3971

```

```
[23]: plot_accuracy(training_2)
```



```
[24]: model_3 = Sequential()

model_3.add(Conv2D(8, (3, 3), activation='relu', input_shape=(256,256,3)))
model_3.add(MaxPool2D((2, 2)))
model_3.add(Flatten())
model_3.add(Dense(32, activation='relu'))
model_3.add(Dense(4, activation='softmax'))
model_3.compile(optimizer='adam',
                loss='categorical_crossentropy',
                metrics=['accuracy'])

model_3.summary()

training_3 = model_3.fit(train_gen,
                        validation_data = test_gen,
                        epochs = 20
                        )
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 254, 254, 8)	224
max_pooling2d_3 (MaxPooling)	(None, 127, 127, 8)	0

2D)

flatten_3 (Flatten)	(None, 129032)	0
dense_6 (Dense)	(None, 32)	4129056
dense_7 (Dense)	(None, 4)	132

=====

Total params: 4,129,412

Trainable params: 4,129,412

Non-trainable params: 0

-----

Epoch 1/20

17/17 [=====] - 22s 1s/step - loss: 5.9153 - accuracy: 0.2744 - val\_loss: 1.8797 - val\_accuracy: 0.3603

Epoch 2/20

17/17 [=====] - 22s 1s/step - loss: 1.3461 - accuracy: 0.3941 - val\_loss: 1.2720 - val\_accuracy: 0.4118

Epoch 3/20

17/17 [=====] - 22s 1s/step - loss: 1.1935 - accuracy: 0.4088 - val\_loss: 1.2593 - val\_accuracy: 0.3456

Epoch 4/20

17/17 [=====] - 22s 1s/step - loss: 1.1040 - accuracy: 0.4365 - val\_loss: 1.2239 - val\_accuracy: 0.4485

Epoch 5/20

17/17 [=====] - 22s 1s/step - loss: 0.9817 - accuracy: 0.5451 - val\_loss: 1.1979 - val\_accuracy: 0.4191

Epoch 6/20

17/17 [=====] - 22s 1s/step - loss: 0.8965 - accuracy: 0.5893 - val\_loss: 1.1076 - val\_accuracy: 0.5221

Epoch 7/20

17/17 [=====] - 22s 1s/step - loss: 0.8289 - accuracy: 0.6611 - val\_loss: 1.1030 - val\_accuracy: 0.5368

Epoch 8/20

17/17 [=====] - 21s 1s/step - loss: 0.7611 - accuracy: 0.7109 - val\_loss: 1.1469 - val\_accuracy: 0.5441

Epoch 9/20

17/17 [=====] - 22s 1s/step - loss: 0.6898 - accuracy: 0.7514 - val\_loss: 1.0856 - val\_accuracy: 0.5662

Epoch 10/20

17/17 [=====] - 22s 1s/step - loss: 0.6354 - accuracy: 0.7864 - val\_loss: 1.1186 - val\_accuracy: 0.5515

Epoch 11/20

17/17 [=====] - 23s 1s/step - loss: 0.5899 - accuracy: 0.8214 - val\_loss: 1.2041 - val\_accuracy: 0.5074

Epoch 12/20


17/17 [=====] - 20s 1s/step - loss: 0.5474 - accuracy:

```

0.8195 - val_loss: 1.1547 - val_accuracy: 0.5368
Epoch 13/20
17/17 [=====] - 22s 1s/step - loss: 0.4869 - accuracy:
0.8471 - val_loss: 1.1122 - val_accuracy: 0.5588
Epoch 14/20
17/17 [=====] - 22s 1s/step - loss: 0.4412 - accuracy:
0.8821 - val_loss: 1.0308 - val_accuracy: 0.5441
Epoch 15/20
17/17 [=====] - 22s 1s/step - loss: 0.4018 - accuracy:
0.8950 - val_loss: 1.0777 - val_accuracy: 0.5662
Epoch 16/20
17/17 [=====] - 22s 1s/step - loss: 0.3533 - accuracy:
0.9061 - val_loss: 1.0271 - val_accuracy: 0.5515
Epoch 17/20
17/17 [=====] - 22s 1s/step - loss: 0.3072 - accuracy:
0.9466 - val_loss: 1.0280 - val_accuracy: 0.5662
Epoch 18/20
17/17 [=====] - 21s 1s/step - loss: 0.2683 - accuracy:
0.9392 - val_loss: 1.1177 - val_accuracy: 0.5588
Epoch 19/20
17/17 [=====] - 28s 2s/step - loss: 0.2318 - accuracy:
0.9650 - val_loss: 1.2484 - val_accuracy: 0.5441
Epoch 20/20
17/17 [=====] - 27s 2s/step - loss: 0.2002 - accuracy:
0.9816 - val_loss: 1.1632 - val_accuracy: 0.5956

```

```
[25]: plot_accuracy(training_3)
```



From the training accuracy plots of the three models we can observe that from epoch 10 the model is starting to over-fit. The training accuracy keeps on increasing, where as the validation accuracy is starting to decrease.