

# dm-1-asst-4

December 16, 2023

```
[61]: import os
import numpy as np, pandas as pd
import keras

import torch
from torch import nn, optim
from torchvision import datasets, models, transforms

import warnings
warnings.filterwarnings('ignore')
```

```
[2]: data_directory = '/content/drive/MyDrive/Dogs_Dataset/images_cropped'
```

```
[3]: true_labels = []
k = 0
for i in os.listdir(data_directory):
    for j in os.listdir(os.path.join(data_directory,i)):
        true_labels.append(k)
    k+=1
true_labels = np.array(true_labels)
```

## Resize

```
[4]: transforms = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
])
```

```
[5]: image_datasets = datasets.ImageFolder(data_directory, transform=transforms)
```

```
[6]: dataloaders = torch.utils.data.DataLoader(image_datasets, batch_size=1,
    ↪shuffle = False)
```

## Normalize

```
[7]: norm_data = []
for i in dataloaders:
    mean,std,var = torch.mean(i[0]),torch.std(i[0]),torch.var(i[0])
```

```
t = (i[0]-mean)/std
norm_data.append(t)
```

## Extract Features

```
[8]: resnet18 = models.resnet18(pretrained=True)
resnet18 = torch.nn.Sequential(*(list(resnet18.children())[:-1]))

resnet18.eval()
```

```
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208:
UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be
removed in the future, please use 'weights' instead.
  warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223:
UserWarning: Arguments other than a weight enum or `None` for 'weights' are
deprecated since 0.13 and may be removed in the future. The current behavior is
equivalent to passing `weights=ResNet18_Weights.IMAGENET1K_V1`. You can also use
`weights=ResNet18_Weights.DEFAULT` to get the most up-to-date weights.
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to
/root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth
100%|          | 44.7M/44.7M [00:00<00:00, 75.4MB/s]
```

```
[8]: Sequential(
  (0): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
  (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (2): ReLU(inplace=True)
  (3): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
ceil_mode=False)
  (4): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
```

```

track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(5): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(6): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,

```

```

track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(7): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)

```

```
)
(8): AdaptiveAvgPool2d(output_size=(1, 1))
)
```

```
[13]: features_extracted = []
      for i in norm_data:
          with torch.no_grad():
              features = resnet18(i)
              features_extracted.append(features)
```

## Dimension Reduction

```
[14]: features_array = []
      for i in features_extracted:
          output_tensor = torch.nn.functional.adaptive_avg_pool2d(i, (1, 1))
          output_tensor = output_tensor.view(512)
          output_tensor = np.array(output_tensor)
          features_array.append(output_tensor)
```

```
[15]: features_array = np.array(features_array)
      features_array.shape
```

```
[15]: (679, 512)
```

## Clustering

### KMeans

```
[47]: from sklearn.cluster import KMeans
      kmeans_1 = KMeans(n_clusters=4, random_state=0, init="random", n_init = 'auto').
          ↪ fit(features_array)
      kmeans_2 = KMeans(n_clusters=4, random_state=0, init="k-means++", n_init = '
          ↪ auto').fit(features_array)
```

```
[48]: from sklearn.cluster import BisectingKMeans
      kmeans_3 = BisectingKMeans(n_clusters=4, init='random').fit(features_array)
```

```
[49]: from sklearn.cluster import SpectralClustering
      spectral = SpectralClustering(n_clusters=4).fit(features_array)
```

### DBSCAN

```
[50]: from sklearn.cluster import DBSCAN

      db = DBSCAN(eps=7, min_samples=2).fit(features_array)
      labels_ = db.labels_
      n_clusters_ = len(set(labels_)) - (1 if -1 in labels_ else 0)
      print("Number of clusters:", n_clusters_)
```

Number of clusters: 4

For  $\text{eps} = 7$  and  $\text{min\_samples} = 2$ ,  $\text{n\_clusters} = 4$

### Agglomerative Clustering

```
[51]: from sklearn.cluster import AgglomerativeClustering
ag_single = AgglomerativeClustering(n_clusters=4, linkage = 'single').
    ↪ fit(features_array)
ag_complete = AgglomerativeClustering(n_clusters=4, linkage = 'complete').
    ↪ fit(features_array)
ag_group = AgglomerativeClustering(n_clusters=4, linkage = 'average').
    ↪ fit(features_array)
ag_ward = AgglomerativeClustering(n_clusters=4, linkage = 'ward').
    ↪ fit(features_array)
```

### Clustering Evaluation

```
[62]: cluster_models = [
    ↪ [kmeans_1, kmeans_2, kmeans_3, spectral, db, ag_single, ag_complete, ag_group, ag_ward]
fms = {'KMeans-Random':0,
      'KMeans-kmeans++':0,
      'Bisecting KMeans':0,
      'Spectral Clustering':0,
      'DBSCAN':0,
      'Agglomerative Clustering - single':0,
      'Agglomerative Clustering - complete':0,
      'Agglomerative Clustering - average':0,
      'Agglomerative Clustering - ward':0}
ss = {'KMeans-Random':0,
      'KMeans-kmeans++':0,
      'Bisecting KMeans':0,
      'Spectral Clustering':0,
      'DBSCAN':0,
      'Agglomerative Clustering - single':0,
      'Agglomerative Clustering - complete':0,
      'Agglomerative Clustering - average':0,
      'Agglomerative Clustering - ward':0}

[63]: from sklearn.metrics import fowlkes_mallows_score
from sklearn.metrics import silhouette_score
for i,j,k in zip(cluster_models,fms,ss):
    predicts = i.fit_predict(features_array)
    fm_score = fowlkes_mallows_score(true_labels,predicts)
    s_score = silhouette_score(np.array(true_labels).reshape(-1,1),predicts)
    fms[j] = fm_score
    ss[k] = s_score
```

Ranking best to worst on fowlkes-mallows\_score

```
[64]: sorted(fms.items(), key=lambda item: item[1])[:, :-1]
```

```
[64]: [('KMeans-Random', 0.8408168530682044),  
      ('Bisecting KMeans', 0.8095209072219244),  
      ('Agglomerative Clustering - ward', 0.7940864447509993),  
      ('Agglomerative Clustering - complete', 0.7748626407369926),  
      ('KMeans-kmeans++', 0.677359269306421),  
      ('Agglomerative Clustering - average', 0.6698996277210605),  
      ('Agglomerative Clustering - single', 0.49711258309478007),  
      ('Spectral Clustering', 0.49711258309478007),  
      ('DBSCAN', 0.4936848010419327)]
```

#### Ranking best to worst on Silhouette score

```
[65]: sorted(ss.items(), key=lambda item: item[1])[:, :-1]
```

```
[65]: [('KMeans-Random', 0.7441528286165865),  
      ('Bisecting KMeans', 0.6890862668884007),  
      ('Agglomerative Clustering - ward', 0.6600248710194033),  
      ('Agglomerative Clustering - complete', 0.6244148818913582),  
      ('Agglomerative Clustering - average', 0.08755115006712988),  
      ('KMeans-kmeans++', -0.12242462372178381),  
      ('Agglomerative Clustering - single', -0.4641194424878871),  
      ('Spectral Clustering', -0.4641194424878871),  
      ('DBSCAN', -0.7414110878570999)]
```